

Identify Type of Disease from Image of Cassava Leaves with Deep Learning Methods

Hanlin Tang, Yuechuan Chen, Hua Tong

Contents

- Background and Motivation
- Methods and Models
- Experiment Details
- Result and Summary



Background and Motivation





Background

Cassava, a major staple food in the developing world that plays a vital role in saving the world from hunger.

At least 80% of household farms in Sub-Saharan Africa grow it.

However, there are viral diseases that preventing good yields.





Diseases



**Cassava Bacterial
Blight (CBB)**



**Cassava Brown Streak
Disease (CBSD)**



**Cassava Green
Mottle (CGM)**



**Cassava Mosaic
Disease (CMD)**



Motivation




- Existing detection methods require the involvement of experts' inspection, which suffers from being labor-intensive, costly, and sometime inaccessible.
- Meanwhile, find out the infected plants sooner could help stopping it from spreading.
- Thus, "identifying disease with image" comes to be an option.

Methods and Models



Models

- Convolutional Neural Network



AlexNet

VGGNet

ResNet

GoogLeNet


ResNeXt

Wide ResNet

Loss function

- Cross-entropy

$$\text{cross-entropy} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^k t_{i,j} \log(p_{i,j})$$

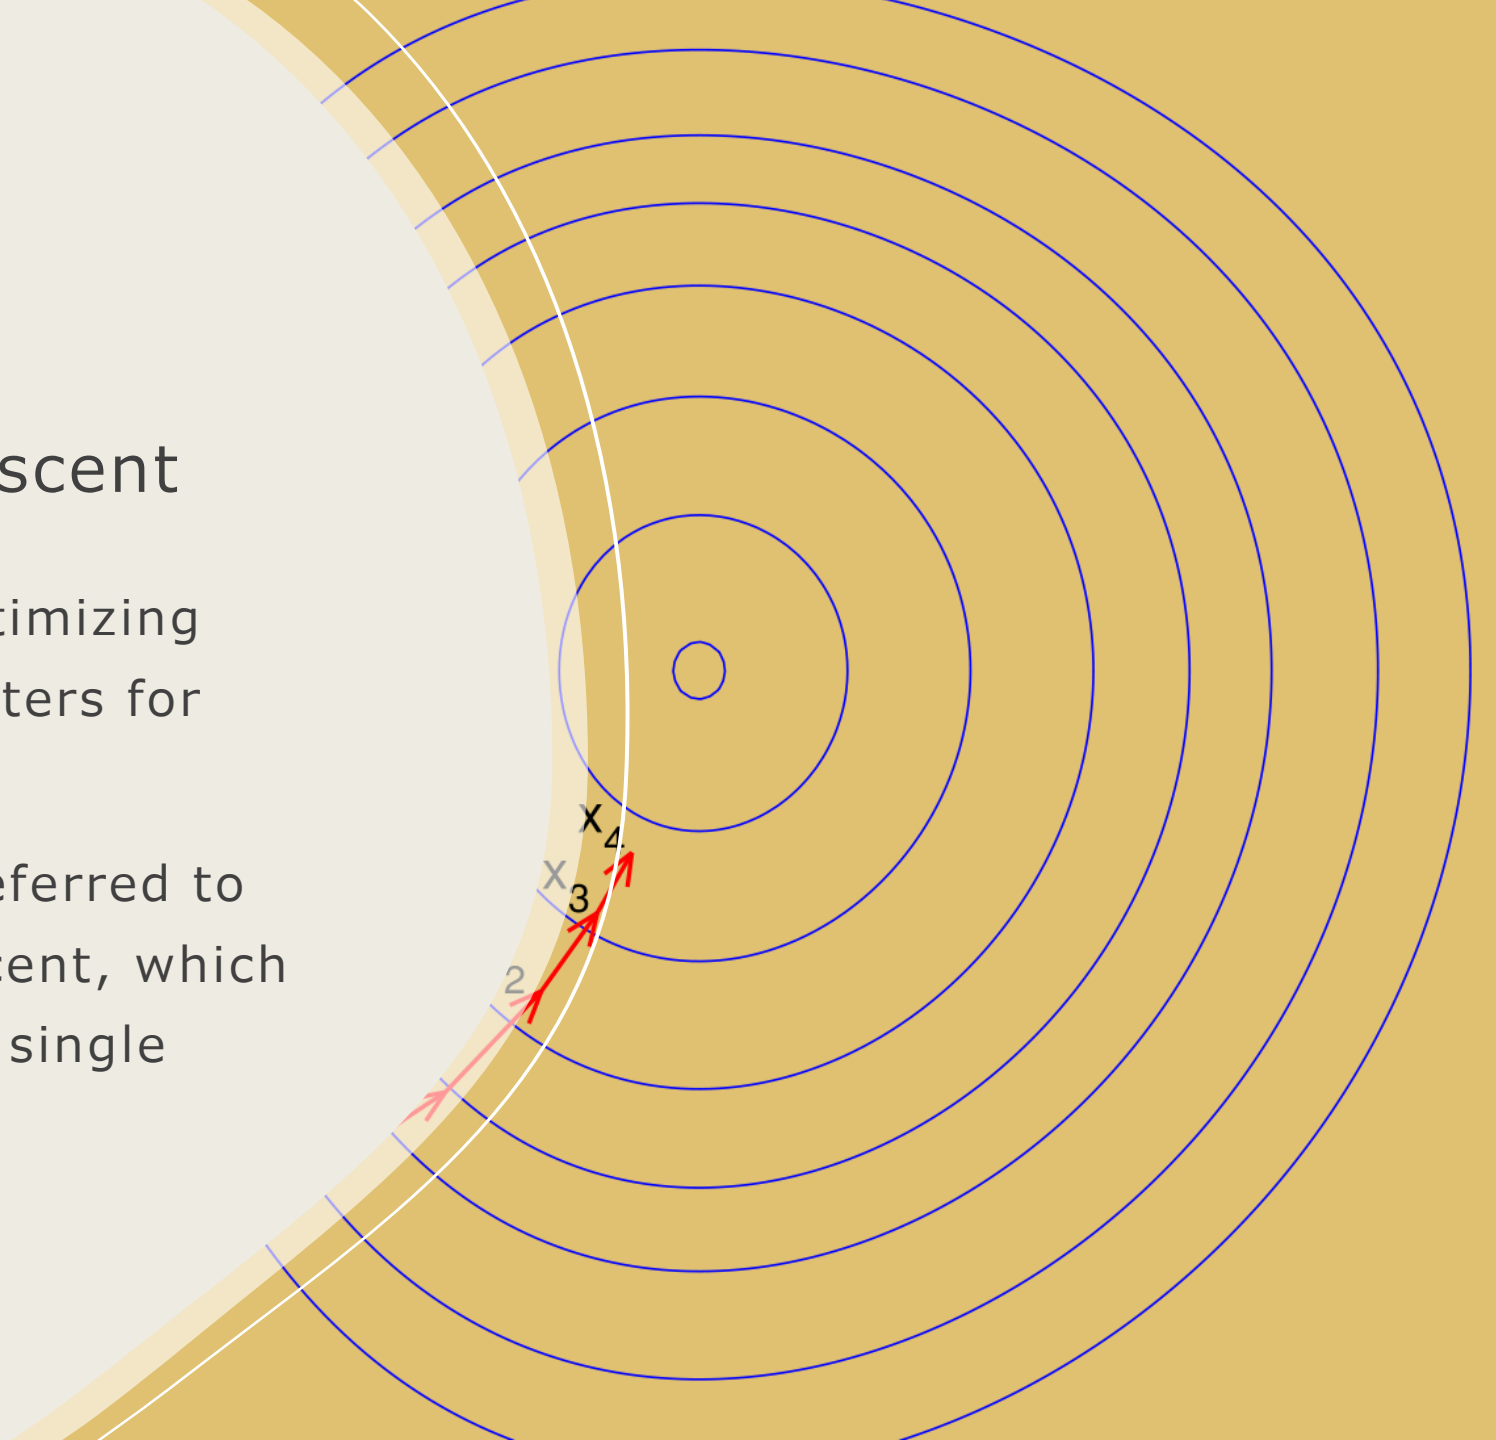


Optimizing

Stochastic gradient descent

Originally referred to the optimizing strategy that update parameters for each single sample.

Nowadays it is more often referred to the mini-batch gradient descent, which take a mini-batch instead of single sample.



Optimizing

Adaptive Moment Estimation (Adam)

Algorithm for first-order gradient-based optimization of stochastic objective functions, based on adaptive estimates of lower-order moments.

The method is straightforward to implement, is computationally efficient, has little memory requirements, is invariant to diagonal rescaling of the gradients, and is well suited for problems that are large in terms of data and/or parameters.

Optimizing

- Early Stopping

Early stopping, as the name stating, is to stop training earlier than our pre-determined epoch number. There are many criteria for doing early stopping, while most of them are based on the validation accuracy.

We designed such an early stopping strategy: “Stop if the max validation accuracy does not increase for 6 epochs, or does not increase for 2 epochs while train accuracy is larger than 99%.”



Experiment Details





Data Set

NUMBERS

- 80% training data (17118)
- 5% validation data (1070)
- 15% test data (3209)
- Around 15000 online test data (not visible)

EACH DATA

- An RGB 3 channel, 800x600 image
- During initial training, we resize to 400x300 for efficiency
- For pretrained models, we normalize with mean [0.485, 0.456, 0.406] and std [0.229, 0.224, 0.225]



Environment

SOFTWARE

- Pytorch on python, including it's embedded models (torchvision.models).



- Overleaf for report, and Powerpoint for slides.

HARDWARE

- One GTX 1650, one RTX 2080, and one RTX 2080 Ti.





General procedure

INITIAL TUNING

- Load model from pytorch package
- Train at most 40 epochs with the early stopping setting
- Generally, use SGD with 0.001 learning rate or 0.0001 learning rate, depends on the result.

FINE-TUNING

- Only tried on Resnet due to lack of resources.
- Make necessary change to the architecture
- Manually adjust learning rate
- Run less epoch each time and keep improving


Result and Summary



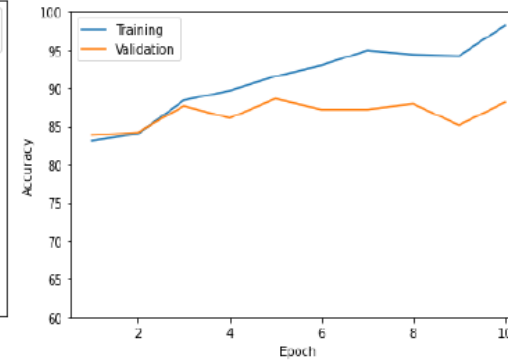
Accuracy

	Training acc	Validation acc	Local test acc	Online test acc
AlexNet	83.07 %	82.43 %	79.37 %	78.21 %
VGG16	98.18 %	88.13 %	85.63 %	84.72 %
Resnet34	99.99 %	84.58 %	82.08 %	80.55 %
GoogLeNet	98.59 %	84.30 %	82.61 %	82.48 %
ResNext	98.26 %	86.73 %	85.76 %	84.37 %
WideResnet	95.86 %	86.17 %	84.11 %	83.78 %


Training Accuracy




AlexNet



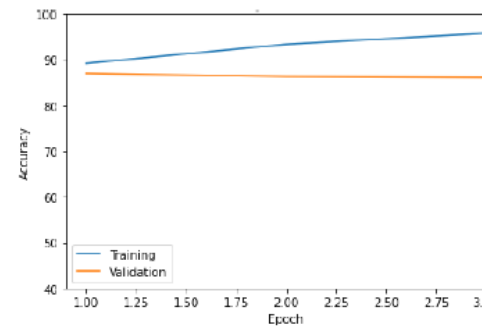
VGGNet




ResNet



GoogLeNet




ResNeXt

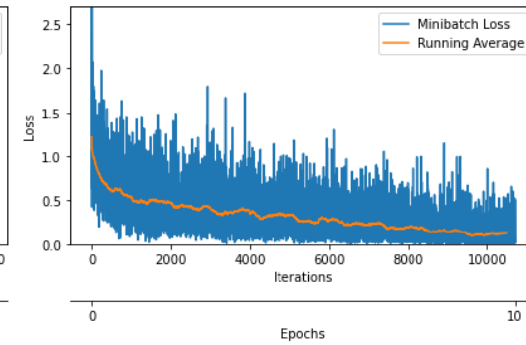


WideResNet


Mini-batch Loss



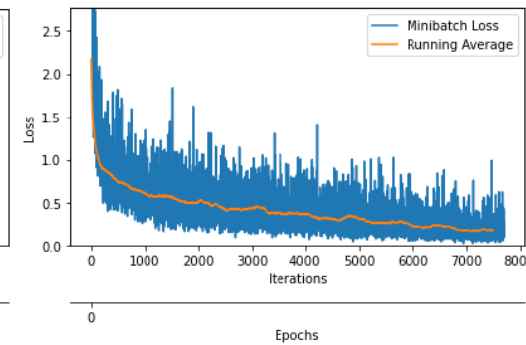
AlexNet




VGGNet



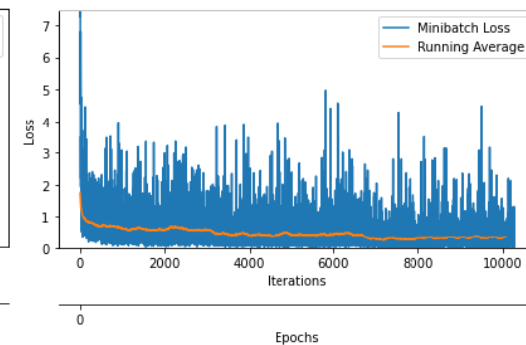
ResNet



GoogLeNet



ResNeXt



WideResNet

Confusion matrix

true label \ predicted label	CBB	CBSD	CGM	CMD	Healthy
CBB	63	23	3	10	35
CBSD	25	181	22	57	44
CGM	3	21	198	106	53
CMD	4	35	45	1876	32
Healthy	25	15	33	71	229

AlexNet

true label \ predicted label	CBB	CBSD	CGM	CMD	Healthy
CBB	72	21	3	10	28
CBSD	23	214	22	30	40
CGM	10	22	232	73	44
CMD	11	37	51	1862	31
Healthy	28	29	17	45	254

ResNet

true label \ predicted label	A	B	C	D	E
A	81	6	3	5	39
B	19	224	16	21	49
C	11	10	281	50	29
D	10	15	36	1874	57
E	28	10	16	27	292

ResNext

true label \ predicted label	CBB	CBSD	CGM	CMD	Healthy
CBB	39	12	1	6	27
CBSD	23	238	7	22	39
CGM	11	14	245	58	53
CMD	9	22	26	1902	33
Healthy	34	10	6	48	275

VGGNet

true label \ predicted label	CBB	CBSD	CGM	CMD	Healthy
CBB	31	19	1	8	42
CBSD	28	248	10	27	33
CGM	12	18	214	60	48
CMD	9	23	41	1827	55
Healthy	42	37	9	36	271

GoogLeNet

true label \ predicted label	A	B	C	D	E
A	85	9	2	3	35
B	24	242	5	9	49
C	11	19	255	41	55
D	17	35	28	1818	94
E	28	25	6	15	299

Wide ResNet



Fine-tuning model

PROCESS

- Due to the limitation of time and resource, we only managed to fine-tuning one model: Resnet.
- Input resized to 512x512 instead of 400x300
- Start with pretrained model
- Replace the last fc layer (2048,1000) with fc (2048,5).
- Trained with smaller mini-batch (8 instead of 16)
- After each epoch, manually decrease learning rate on acc fluctuating
- Eventually run 6 epochs

RESULT

- Train acc: 91.15 %, validation acc: 88.60 %
- Test acc: 87.19%, online test acc: 86.75 %

true label	CBB	CBSD	CGM	CMD	Healthy	
	CBB	72	21	3	10	28
	CBSD	23	214	22	30	40
	CGM	10	22	232	73	44
	CMD	11	37	51	1862	31
	Healthy	28	29	17	45	254
		predicted label				
		CBB	CBSD	CGM	CMD	Healthy



Summary

We practiced our skills in deep learning and meet the target we set at the beginning of the project.

We also learned the complexity of tuning and how cost it could be to run deep learning training.

The final result, though not outstanding, is very much usable and helpful.

Yet, more fine-tuning is expected if time and resource allowed.

Thank You

Hanlin Tang, htang79@wisc.edu

Yuechuan Chen, ychen959@wisc.edu

Hua Tong, htong24@wisc.edu

Github: <https://github.com/Tonghua35/453-Final-Project>